

Analysis of iSCSI Target Software

FUJITA Tomonori

NTT Cyber Solutions Laboratories

tomof@acm.org

SNAPI'04, Antibes Juan-les-Pins, France

What is iSCSI

SCSI over TCP/IP


- SCSI commands are encapsulated into IP packets
- A host and a storage system are connected with Ethernet

New networked storage technology

- Today the dominant networked storage technology is Fibre Channel

Advantages

- IP networks infrastructure is inexpensive
- We are familiar with IP network and SCSI technology



Various iSCSI storage systems

Commercial solutions

- Optimized operating systems and specialized hardware
 - ▷ *iSCSI NIC, NVRAM, etc*
- Many vendors sell iSCSI storage systems
 - ▷ *IBM, EMC, Hitachi, NetApp, etc*

Open source solutions

- Linux and commodity hardware
 - ▷ *PC and Gigabit NIC*
- Two well-known open source iSCSI targets
 - ▷ *UNH - New Hampshire Univ. and HP*
 - ▷ *Ardis - Ardis Technologies Corp.*

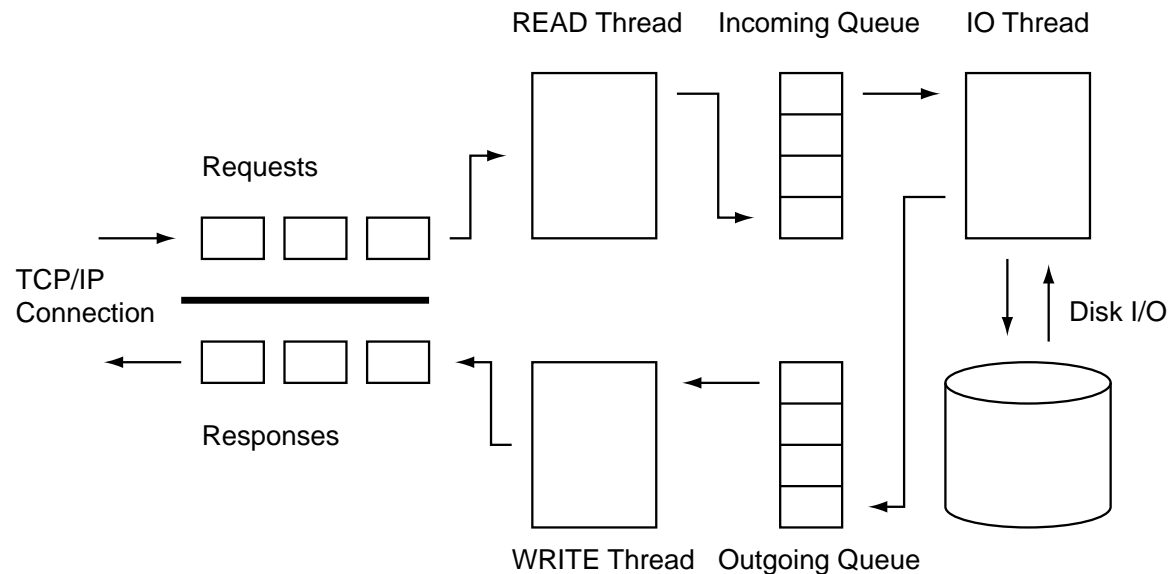
Questions

What are the differences among open source targets?

- UNH-disk
 - ▷ *UNH iSCSI target configured as DISKIO mode*
- UNH-file
 - ▷ *UNH iSCSI target configured as FILEIO mode*
- Ardis
- Threaded-Ardis
 - ▷ *Our new implementation based on the Ardis code*

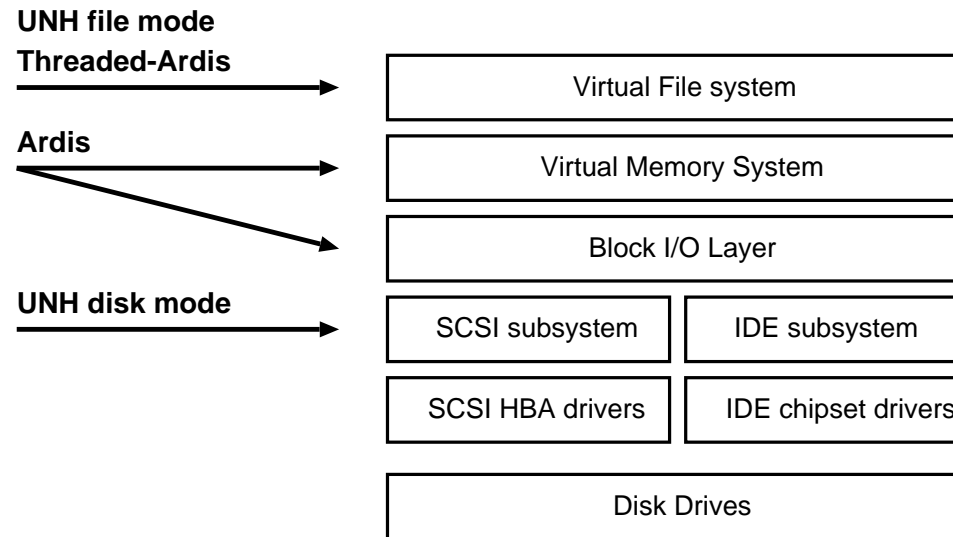
Are open source solutions comparable to commercial products?

Architecture overview



- Three kernel threads provides the major part of the iSCSI target functionality
 - ▷ *Read thread receives data from initiators*
 - ▷ *I/O thread performs I/O operations*
 - ▷ *Write thread sends responses to initiators*
- There are major differences in I/O thread design

I/O thread designs



Linux kernel provides several interfaces for I/O operations

- Virtual file system (vfs) interface
 - ▷ *UNH-file and Threaded-Ardis*
- Modified Virtual Memory subsystem and block I/O layer
 - ▷ *Ardis*
- SCSI subsystem interface
 - ▷ *UNH-disk*

Assessment criteria - kernel modifications

Ardis modifies the source code of Linux kernels

- Ardis implements own I/O functions by using a modified kernel
 - ▷ *Improved performance because Linux standard functions are not optimized for iSCSI targets*
- **More complicated code**
 - ▷ *Increased cost for implementation*
 - ▷ *Increased cost for maintenance of the code*
 - ▷ *Lowered stability*

Assessment criteria - Interoperability

UNH-file breaks interoperability with direct-attached storage

- UNH-file sends the response of WRITE commands before the data are written to disk
 - ▷ *Improved performance due to delayed disk I/O operations*
- **It require software changes in the existing operating system or application**
 - ▷ *Data integrity protection techniques like journaling don't work*
 - ▷ *Data corruption may happen after a system crash*

Delayed write is different from disk drive cache

- Delayed write aggressively reorders write operations
 - ▷ *It breaks data integrity protection techniques*

Assessment criteria - Disk management

UNH-disk suffers poor disk management features

- UNH-disk design can use only SCSI disk drives
 - ▷ *The SCSI subsystem is lower than the block I/O layer providing disk management features*
- UNH-disk doesn't supports various block devices
 - ▷ *IDE*
 - ▷ *Serial-ATA*
- UNH-disk doesn't supports virtual block devices
 - ▷ *Flexible storage management (LVM)*
 - ▷ *Redundancy (software RAID)*

Assessment criteria - Performance

UNH-disk suffers poor read performance

- UNH-disk design cannot use page cache minimizing disk I/O by storing data
 - ▷ *The SCSI subsystem is lower than VM system providing page cache functionality*
- Every SCSI read command invokes a disk I/O
 - ▷ *It leads to poor read performance*

UNH-file suffers poor performance

- UNH-file design cannot perform SCSI commands simultaneously
 - ▷ *The vfs interface works synchronously - The I/O thread sleeps until I/O completions*
- All I/O operations are serialized unnecessarily
 - ▷ *It leads to poor performance*
- Threaded-Ardis uses multiple I/O threads to avoid this problem
 - ▷ *This design enables Threaded-Ardis to handle SCSI commands synchronously*

System Comparison

Examined targets

- Ardis
- Threaded-Ardis
- UNH-disk
- UNH-file
- UNH-file-sync
 - ▷ *Modified UNH-file providing interoperability about WRITE commands*
- Entry-class commercial iSCSI target system
 - ▷ *The product details are confidential*

Benchmark software

- Microbenchmarks
 - ▷ *Sequential write*
 - ▷ *Sequential read*
- Macrobenchmarks
 - ▷ *Postmark - write intensive*
 - ▷ *Network server benchmark - read intensive*

Experimental infrastructure

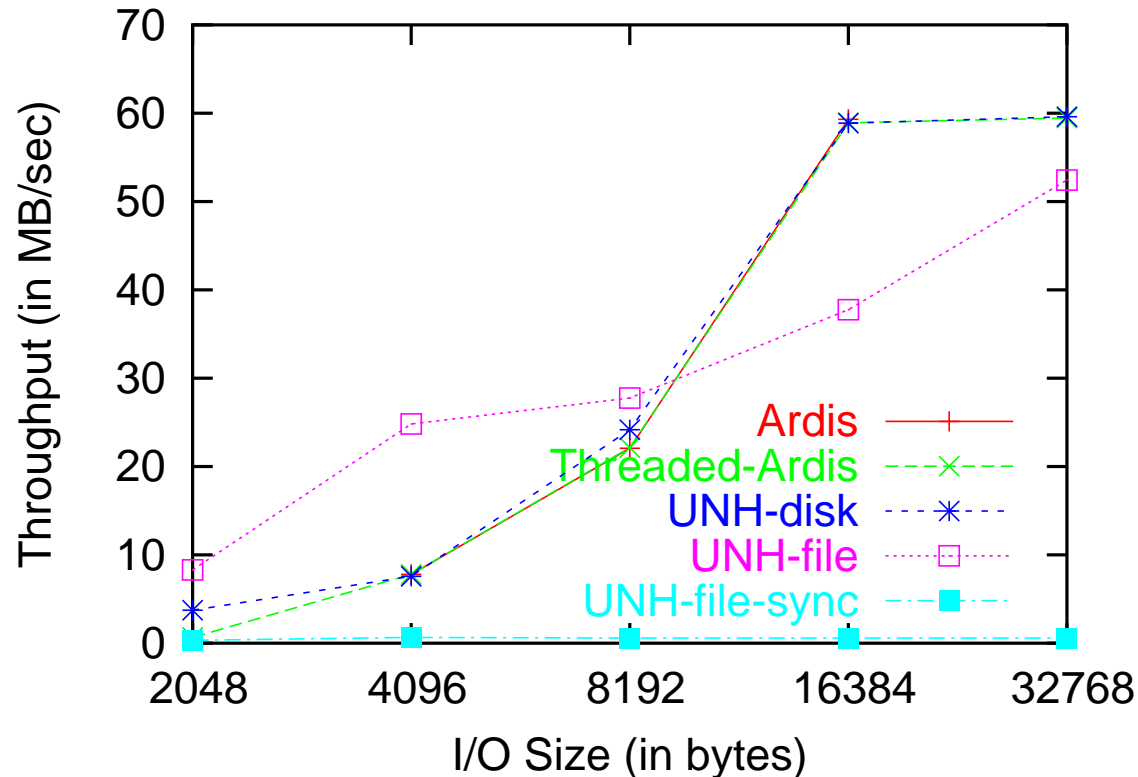
Initiator

- Linux kernel 2.6.4
- 2 GHz Xeon processor
- 1 GB main memory
- Cisco initiator 4.0.1.1

Target


- Linux kernel 2.4.25
- 2 GHz Xeon processor
- 2 GB main memory
- Maxtor Atlas 10K, 36.7 GB 10,000 RPM SCSI disks
- LSI Logic 53C1030 Ultra320 SCSI chip

Microbenchmarks - Write




- UNH-file produces better performance due to delayed write
 - ▷ *No I/O operations are performed up to 8 KB*
- No clear gain from specialized functions
- Ardis couldn't complete the benchmarks with 2KB or 32KB due to its bugs.

Microbenchmarks - Read




- Ardis performance is slightly better than the others
- UNH-file performance is poor
 - ▷ *It needs to perform I/O operations*
 - ▷ *It cannot handle them simultaneously*

Postmark



- Postmark is designed to measure performance in the ephemeral small-file workloads seen by ISP
 - ▷ *20,000 files, 50,000 transactions, and file sizes of between 512 bytes and 16 KB*
- Comparable throughputs except for the UNH-file-sync and UNH-file targets

Network Server benchmark



- The benchmark repeatedly reads many files in one directory in random order
 - ▷ 512, 2 MB files 4,096 times in total
 - ▷ 1,024, 2 MB files 4,096 times in total
- Large effects of page cache on the performances
 - ▷ Page cache hit rate of 85.4% and 25.4% respectively at the target

Summary

Our design advantages

- Low cost
 - ▷ *No kernel modification*
 - ▷ *No changes to existing software*
- Rich disk management features
 - ▷ *Exploiting kernel virtualization of block devices*
- Comparable performances in common workloads
 - ▷ *The iSCSI target modifying a Linux kernel*
 - ▷ *Entry-class commercial iSCSI target system*

Current development status

- The code was released publicly at May 25, 2004
 - ▷ <http://sourceforge.net/projects/iscsitarget/>
- The code has been actively maintained by several developers